

Chapter 5: Calculating Network Error with Loss

To train a model, we tweak the weights and biases to improve the model's accuracy and confidence. To do this, we calculate how many errors the model has.

The **loss function**, also referred to as the **cost function** algorithm that quantifies how wrong a model is. Loss is the measure of the loss function. And, we ideally want it to be 0.

You may wonder why we do not calculate the error of a model based on the argmax accuracy. Recall our earlier example of confidence: [0.22, 0.6, 0.18] vs [0.32, 0.36, 0.32]. If the correct class were indeed the middle one (index 1), the model accuracy would be identical between the two above. But are these two examples really as accurate as each other? They are not, because accuracy is simply applying an argmax to the output to find the index of the biggest value. The output of a neural network is actually confidence, and more confidence in the correct answer is better.

Loss function for:

Linear Regression: Squared error or mean squared error

Binary Logistic Regression: Log Loss

Classification: Categorical Cross-Entropy Loss

Cross-entropy is a loss function that measures the difference between two probability distributions:

- The **true distribution** (ground truth labels).
- The **predicted distribution** (model outputs, typically after softmax or sigmoid).

Special cases of cross entropy:

Categorical cross-entropy is a special case for **multi-class classification**,

Log Loss is another special case for two classes i.e **binary classification**.

Categorical Cross-Entropy Loss

As we are classifying the data, we're using this loss function.

Categorical cross-entropy is explicitly used to compare a “ground-truth” probability (y or “targets”) and some predicted distribution (y-hat or “predictions”).

The formula for calculating the categorical cross-entropy of y (actual/desired distribution) and y-hat (predicted distribution) is:

$$L_i = - \sum_j y_{i,j} \log(\hat{y}_{i,j})$$

Where L_i denotes sample loss value, i is the i -th sample in the set, j is the label/output index, y denotes the target values, and \hat{y} denotes the predicted values.

The softmax activation function returns a probability distribution over all of the outputs. Cross-entropy compares two probability distributions.

Let's say:

softmax_output = [0.7, 0.1, 0.2] ; which is the predicted probability distribution

targeted probability distribution = [1, 0, 0]

*Arrays or vectors like this are called **one-hot**, meaning one of the values is “hot” (on), with a value of 1, and the rest are “cold” (off), with values of 0. When comparing the model's results to a one-hot vector using cross-entropy, the other parts of the equation zero out, and the target probability's log loss is multiplied by 1, making the cross-entropy calculation relatively simple. This is also a special case of the cross-entropy calculation, called **categorical cross-entropy**.*

So,

$$\begin{aligned} L_i &= - \sum_j y_{i,j} \log(\hat{y}_{i,j}) = -(1 \cdot \log(0.7) + 0 \cdot \log(0.1) + 0 \cdot \log(0.2)) \\ &= -(-0.35667494393873245 + 0 + 0) = 0.35667494393873245 \end{aligned}$$

While programming,

```
loss = -(math.log(softmax_output[0])*target_output[0] +
          math.log(softmax_output[1])*target_output[1] +
          math.log(softmax_output[2])*target_output[2])
```

As, values other than the target value in the target output is zero and the target value itself is 1 so the loss function can be simplified as

```
loss = -math.log(softmax_output[0])
```

Hence, the formula

$$L_i = - \sum_j y_{i,j} \log(\hat{y}_{i,j})$$

Is simplified to:

$$L_i = -\log(\hat{y}_{i,k}) \quad \text{where } k \text{ is an index of "true" probability}$$

Where L_i denotes sample loss value, i is the i -th sample in a set, k is the index of the target label (ground-true label), y denotes the target values and \hat{y} denotes the predicted values.

What was once an involved formula reduces to the negative log of the target class' confidence score

the example confidence level might look like [0.22, 0.6, 0.18] or [0.32, 0.36, 0.32]. In both cases, the argmax of these vectors will return the second class as the prediction, but the model's confidence about these predictions is high only for one of them. The Categorical Cross-Entropy Loss accounts for that and outputs a larger loss the lower the confidence is:

```
import math

print(math.log(1.))
print(math.log(0.95))
print(math.log(0.9))
print(math.log(0.8))
print('...')
print(math.log(0.2))
print(math.log(0.1))
print(math.log(0.05))
print(math.log(0.01))
```

```
>>>
0.0
-0.05129329438755058
-0.10536051565782628
-0.2231435513142097
...
-1.6094379124341003
-2.3025850929940455
-2.995732273553991
-4.605170185988091
```

When the confidence level equals 1, meaning the model is 100% “sure” about its prediction, the loss value for this sample equals 0. The loss value raises with the confidence level, approaching 0.

Targets can be either:

- **one-hot encoded**, where all values, except for one, are zeros, and the correct label’s position is filled with 1.

Eg:

```
class_targets = np.array([[1, 0, 0],
                          [0, 1, 0],
                          [0, 1, 0]])
```

- **Sparse**, which means that the numbers they contain are the correct class numbers.

Eg:

```
class_targets = [0, 1, 1]
```

The check can be performed by counting the dimensions — if targets are single-dimensional (like a list), they are sparse, but if there are 2 dimensions (like a list of lists), then there is a set of one-hot encoded vectors.

If sparse target class:

```
# Probabilities for target values -  
# only if categorical labels  
if len(class_targets.shape) == 1:  
    correct_confidences = softmax_outputs[  
        range(len(softmax_outputs)),  
        class_targets  
    ]
```

If one hot encoded target class:

```
# Mask values - only for one-hot encoded labels  
elif len(class_targets.shape) == 2:  
    correct_confidences = np.sum(  
        softmax_outputs * class_targets,  
        axis=1  
    )  
  
softmax_outputs = [[0.7, 0.1, 0.2],  
                  [0.1, 0.5, 0.4],  
                  [0.02, 0.9, 0.08]]  
  
class_targets = [0, 1, 1]  
  
for targ_idx, distribution in zip(class_targets, softmax_outputs):  
    print(distribution[targ_idx])
```

```
softmax_outputs = np.array([[0.7, 0.1, 0.2],
                             [0.1, 0.5, 0.4],
                             [0.02, 0.9, 0.08]])
class_targets = [0, 1, 1]

print(softmax_outputs[[0, 1, 2], class_targets])
```

```
print(softmax_outputs[
    range(len(softmax_outputs)), class_targets
])

print(-np.log(softmax_outputs[
    range(len(softmax_outputs)), class_targets
])))
```

If confidence score is 0

Now when confidence score is calculated, there arises another problem, what if the model is absolutely wrong and has a confidence score of 0, Then $-\text{np.log}(0)$ yields negative infinity, which is a problem for further calculation ie to calculate average loss or the gradient later on.

To fix it:

a. Adding a small value to the confidence score for example $1e-7$:

```
-np.log(1e-7)
```

```
>>>
16.11809565095832
```

This will insignificantly impact the result but yields additional two issues;

- When the model is fully correct ie confidence score is 1, loss becomes a negative value instead of being 0.

```
-np.log(1+1e-7)
```

```
>>>
-9.999999505838704e-08
```

- Shifting confidence towards 1, even if by a very small value.

b. Clip values from both sides by the same number, $1e-7$ in our case:

The lowest possible value will become $1e-7$ (like in the demonstration we just performed) but the highest possible value, instead of being $1+1e-7$, will become $1-1e-7$ (so slightly less than 1).

This will prevent loss from being exactly 0, making it a very small value instead, but won't make it a negative value and won't bias overall loss towards 1.

```
y_pred_clipped = np.clip(y_pred, 1e-7, 1 - 1e-7)
```

Doing so, the lowest value of y_{pred} cannot be less than $1e-7$ and the highest cannot be more than $1 - 1e-7$, if exists it gets clipped to these threshold values.

Eg.
arr = np.array([1, 5, 10, 15])
clipped = np.clip(arr, 3, 12)
print(clipped)

```
>>>  
[ 3  5 10 12]
```

Accuracy, which describes how often the largest confidence is the correct class in terms of a fraction